

Voyager: Critical Additions & Strengthening Recommendations

MUST ADD

1. Progressive Complexity Model

Problem: Full architecture (memo → parse → critique → quant → execution → monitoring) risks overwhelming users.

Solution:

- **Tier 1 (Day 1):** Simple thesis entry + basic performance tracking + daily alerts
- **Tier 2 (Week 2):** Full quant evaluation + LLM critique
- **Tier 3 (Month 1):** Cross-thesis analysis + automated rebalancing

Default to "light-touch monitoring" as your Phase 1 Q&A emphasizes—reserve heavy analytics for thesis birth, review, and death.

2. Cold Start Content

Problem: Empty system has no value on day one.

Solution:

- Pre-loaded demo portfolios ("Gold vs Inflation," "Tech vs Value") users can explore
- One-click portfolio import from brokerages to populate context immediately
- Template library for common thesis types (carry trade, sector rotation, regime shift)

3. Thesis Versioning System

Problem: Reality drifts from original thesis—positions added for new reasons, macro overlays change sizing.

Solution:

- When substantial changes occur, create "Thesis v2" referencing v1
- Preserve analytical history: "This position started as Thesis A (inflation hedge) but you added 30% more for Thesis B (USD weakness)"
- Your architecture already contemplates this—make it explicit in the data model

4. Collaboration Scaffolding (Even for Solo Users)

Problem: Even solo investors work with advisors/partners.

Solution (minimal):

- Thesis sharing (read-only links)
- Comment threads on observations
- "Compare my thesis with..." feature

This expands addressable market without architectural rewrites.

MUST EMPHASIZE

1. The "Aha" Moment Hierarchy

Your architecture has multiple value propositions—prioritize:

PRIMARY (the moat):

- Automated disconfirmer alerts + AI summaries connecting market moves to thesis reasoning
- This is "having an analyst constantly checking if your assumptions still hold"—nobody else does this well

SECONDARY:

- Cross-thesis coherence checking (surfaces hidden concentration risk)
- Conviction-to-sizing translation (quant-informed position sizing)

Lead with disconfirmation monitoring in all messaging. The others are features; this is the core insight.

2. The Daily Briefing as Habit Engine

Your Phase 1 Q&A vision—market dashboard + narrative headlines + PnL attribution + macro radar—is the engagement hook.

Polish this ruthlessly. If opening Voyager becomes the fastest way to answer "what happened overnight and does it matter to my book?", you've created a daily habit. Habits create retention. Retention creates willingness to pay.

Make this the hero feature in demos, onboarding, and marketing.

3. Provenance & User Control

Your LLM architecture includes hallucination controls, but users will still distrust AI initially.

Always show provenance:

- "This risk flag appeared because Observation X mentioned Y"
- "This critique is based on: [list sources]"
- Let users edit/override any LLM output

Your design says "suggest but never auto-mutate"—emphasize this as a design philosophy. Trust builds when users feel in control, not automated.

MUST STRENGTHEN

1. Thesis Entry UX

Current: Memo → parse → critique → quantify is conceptually elegant but potentially clunky.

Strengthen with:

- Guided thesis builder with prompts for each structured field
- LLM auto-extraction from uploaded memos (user confirms, doesn't retype)
- Smart defaults: "You mentioned 'gold' and 'inflation'—suggest GLD 80%, TIP 20%?"
- In-line critique as user types: "Your hypothesis implies USD weakness, but your expression is USD-neutral—intentional?"

The thesis object is your foundation. Make creating it feel effortless, not like filling out a tax form.

2. "Why My Portfolio Moved" as Core Query

Your daily update session generates summaries, but the user need is specific: **"Why did I make/lose money today, and does it validate or contradict my theses?"**

Strengthen:

- Make this a one-click action ("Explain Today")
- Format as: PnL breakdown by thesis → which thesis drove gains/losses → what market moves caused this → do those moves confirm or contradict your original reasoning
- This is the "synthesis layer" your Phase 1 Q&A reserves for humans—but the system should organize the raw materials perfectly

3. Risk Transparency in Quant Layer

Your E2 architecture computes metrics (Sharpe, max drawdown, VaR) but abstracts complexity.

Strengthen:

- Show assumptions explicitly: "This backtest assumes daily rebalancing—your actual execution will differ"
- Highlight edge cases: "This thesis has negative equity in backtest after 2022 (short position dominated)—be aware"
- Offer scenario sensitivity: "If inflation rises 2% more than expected, your thesis gains X% but concentration risk increases Y%"

Sophisticated users will trust the system more if it shows its work and limitations.

4. Error Handling as Feature

Your architecture has robust error handling (404s, LLM failures, missing data), but treats these as edge cases.

Reframe as features:

- When LLM returns `insufficient_context: true`, show: "I don't have enough data to critique this thesis yet. Add 3 observations or wait 5 trading days for price history."
- When quant evaluation fails due to missing prices, suggest: "Asset XYZ lacks historical data. Use proxy ABC instead?"
- When scenario analysis shows extreme sensitivity, flag it: "Warning: This thesis is highly sensitive to interest rate moves—consider hedging with..."

Errors are teaching moments. Use them to make users better investors.

VALIDATE IMMEDIATELY

Before building more, test these three hypotheses:

1. Daily Habit Test

Show 10 target users a prototype of the morning briefing. Ask: "Would you open this every weekday morning?"

If <7 say yes, the engagement hook isn't strong enough.

2. Thesis Structure Test

Have 5 users articulate a real investment idea using your structured format (hypothesis, drivers, disconfirmers, expression).

If they struggle or feel constrained, the entry UX needs work. If they find it clarifying, you've validated the core insight.

3. Continuous Monitoring Value Test

Show users a simulated "thesis health check" after 30 days: driver changes, disconfirmer alerts, performance vs. expectation. Ask: "Is this valuable enough to pay for?"

If they say "interesting but I'd check manually," the automation isn't compelling enough. If they say "I'd pay to not miss this," you've found product-market fit.

ONE-SENTENCE STRENGTHENING IMPERATIVE

Make the daily briefing so good that opening Voyager becomes the fastest way to answer "what moved and does it matter to my theses?", then nail the thesis entry UX so creating structured investment ideas feels effortless, not bureaucratic.