QUESTION 1

When I open Slice on a weekday morning, I want to be dropped immediately into the world itself — the cross-asset surface that every PM, trader, and macro thinker wakes up to. The first thing on the screen is a clean, high-signal **macro/market dashboard**: equities (SPY, QQQ, Nikkei, EuroStoxx), long-end and short-end rates, real yields, breakevens, gold, oil, DXY, and a few G10 pairs. No stories yet — just the state of the battlefield. I see what moved, what didn't, and whether the overnight tape was coherent or strange. This is the opening chord that tunes my intuition.

Beneath it sits the **market headlines module**, still within the same "Morning Briefing" card — a narrative layer that gives contour to the numbers. Equity headlines, rates headlines, FX commentary, commodities color. A mix of Reuters, FT, and curated macro signals. Not noise — the narrative structure that tells me *why* things are behaving the way they are, or whether they are behaving without explanation at all.

Below that, I fall into my own reflection: **PnL + attribution**, shown cleanly and without drama. How my book interacted with the world that just passed. What drove gains, what chipped away, what behaved out of character. It's not about judgment — it's about staying awake to the shape of my exposures.

Finally, wrapped inside the same Morning Briefing card, are the three blocks that I consider indispensable and non-negotiable: **Rates, FX, and the Macro Event Radar**. They're compact: 2s/10s, reals, breakevens; DXY plus a couple crosses; today's top economic prints with expected vs prior. These aren't "extras." They are the scaffolding. Together they tell me whether today is likely orderly or dangerous, whether positioning matters or flows matter, whether regime signals are shifting under the surface.

QUESTION 2

When I bring a macro view into Slice, it does not start as a prompt or a bullet point — it begins as a *memo*. A small essay, built in collaboration with my team and with GPT as an assistant:

charts assembled, cross-asset relationships diagrammed, historical episodes referenced, pricing trends highlighted, regime markers identified. This memo is the birthplace of the thesis. It is where the intellectual architecture is written down in full — the narrative spine, the analytics, the intuition.

Once I drop that memo into Slice, the machine takes its turn. Slice parses it first — extracting the hypothesis, the drivers, the disconfirmers, the expression, the dates, the proxies. Then it critiques it, asking the questions that force me to sharpen causal pathways and confront hidden assumptions. And then it quantifies it — pulling historical analogs, linking the thesis to real data, mapping the expression to proxies, and checking whether the structure of the idea has any statistical backbone at all. This is the order: **parse → critique → quantify**. The human writes the stone; the system chisels it into form.

From that point forward, my daily observations are not bricks thrown at the thesis nor automatic edits to its structure. Rather, Slice uses my observations — structured, embedded, retrieved — to *suggest* updates. It highlights contradictions, drift, confirmations, risks, or disconfirming data, but it never overwrites the thesis without my approval. The thesis is a structured object and a living narrative, but it is not an auto-mutating one. Slice interrogates; I decide.

QUESTION 3

When I sit with a thesis inside Slice, the minimum quantitative output I expect is the backbone — the historical behavior of the expression. Before I see a single chart of sensitivities or a single scenario matrix, I need the history: how the trade has behaved through cycles, through different real-yield regimes, through liquidity shocks, through cross-asset dislocations. This is the foundation. The driver state — where the animating variables stand today, and how they've moved since the thesis was born — is the animus. And the live performance of the thesis expression is the counter: the sober reminder of how the world is actually treating my idea in real time.

On a day-to-day basis, Slice should not unload the full quant armory on me. Full analytics — regressions, factor models, full backtests, VaR, and scenario sweeps — are for three moments only: the birth of the thesis, its periodic formal review, and its burial. In the daily rhythm, I want light-touch monitoring that tells me whether the drivers are drifting toward or away from the thesis, whether any disconfirmers have flashed, and whether the portfolio is quietly leaning into

or against the view. When needed, I can drill down into rolling correlations or beta shifts, but these live in the second layer, not the surface.

QUESTION 4

The end of a Slice session depends on the kind of session I'm in. If I'm simply checking on the book — taking in the morning briefing, seeing how the exposures, drivers, and thesis expression are behaving — then the purpose is not execution. The purpose is provocation. Slice should ask the questions that spark reflection, nudge me toward awareness of drift, highlight conflicts I didn't see, or show me where the numbers whisper something meaningful. Macro PMs are not constantly firing trades; they are constantly thinking. The system should support that. The "end" of this kind of session is a thought — not a ticket.

But when I'm in a thesis-centered trade-building session, the ending shifts. Here the system should converge toward a proposal: a sizing recommendation, a structure, a way to express the thesis through GLDM or RINF or rates or FX, tied to my risk rails and exposures. I don't want a single deterministic answer; I want a set of options — a partially automated, highly filtered set of scenarios, with Slice explaining the trade-offs and honing in, through iteration, on the best option. I choose the path, but the machine narrows the corridor.

Outside either of these sessions — at the pure portfolio level — Slice operates with autonomy. It should rebalance within its rails on its own, adjusting small drifts, maintaining risk posture, and executing pre-declared tactical actions like stops, take-profit levels, or timed adds that were explicitly encoded inside the thesis expression. These are not new ideas; they are the mechanical expressions of the idea already approved. The human defines the structure; the system keeps it coherent.

Slice does not pretend to be a trader. It is the chisel that keeps the stone aligned — and when the human chooses to sculpt, it brings the tools to the table.

QUESTION 5

When I open Slice the next morning, I want it to remember what *matters* — the shifts in drivers, the movement in disconfirmers, how the thesis expression behaved, where the exposures drifted. But I don't want Slice to explain the world to me. The point is not to be told *why* things happened; the point is for the machine to surface the structural changes so that **I can see the pattern myself**. That is the human domain. Slice lays out the threads — I trace the tapestry.

It should bring forward the observations I wrote the day before, not as interpretations but as reminders. It should show how those intuitions aged relative to the data, what held up, what didn't. It should highlight the tension points — places where the world nudged the thesis or contradicted it — without collapsing them into a narrative. Interpretation is the human craft. The machine simply ensures nothing is forgotten.

On the portfolio side, it should tell me what small autonomous actions it took — rebalances, triggered stops or adds — so I can see how the structure of the book interacted with the structure of the idea. Again, no storytelling, no causal gloss. Just the *facts* of what the system did and what moved.

Slice isn't meant to hand me a conclusion. When I return each morning, it should present continuity, context, and cleanly surfaced deltas — and leave the synthesis to me. The human sees across trends; the human makes the bigger picture. Slice is the instrument panel, not the pilot.