
PHASE 7 COMPLETION REPORT (DETAILED & ROBUST)

Executive Summary

Phase 7 has been fully completed and validated. All intended capabilities, engines, routes, contexts, and orchestration integrations are in place, functioning as designed, and backed by unit + workflow tests. The strategic intelligence layer is now structurally integrated and stable.

1. Completed Functional Deliverables

A. New Strategic Intelligence Engines (All Implemented)

Each engine is fully implemented and callable through the public API:

1. **Long-Horizon Reasoning Engine**
 - Performs multi-month horizon reasoning over a primary thesis.
 - Builds deterministic long-horizon context from thesis + macro snapshot.
 - Declares analytical structure within user_text before orchestration.
2. **Strategy Recommendation Engine**
 - Ingests all active theses + risk profile.
 - Builds strategy context deterministically.
 - Provides structured instructions for strategic positioning.
3. **Portfolio Diagnostics Engine**
 - Ingests risk snapshot and stubbed portfolio representation.
 - Performs exposure diagnostics and risk concentration assessment.
 - Guarantees deterministic context build.
4. **Narrative Coherence Engine**
 - Ingests complete thesis set + macro snapshot.
 - Builds narrative linkage context.
 - Produces structured macro-portfolio narrative guidance.

All four engines:

- call ContextBuilder correctly,
 - serialize context correctly into user_text,
 - invoke orchestrator using correct Phase-6 contract.
-

2. Architectural Compliance

A. Context Builder

- Phase-7-specific deterministic builders implemented.
- No LLM calls.
- No mutation.
- Data-access only.
- Clear “kind” markers for each context type:
 - long_horizon_context
 - strategy_context
 - portfolio_diagnostics_context
 - narrative_coherence_context

B. Orchestrator Boundary Integrity

- Zero changes to OrchestratorClient public API.
- No unapproved params introduced.
- engines only use:
 - run_standard
 - run_analyst
 - run_concise
- All engines serialize context into user_text (correct behavior).

C. API Layer

Phase 6 API preserved exactly.

No regression.

Phase 7 adds:

- /intel/horizon
- /intel/strategy
- /intel/diagnostics
- /intel/narrative

All are thin delegations, purity maintained.

3. Testing Coverage Status

A. Legacy Tests

100% pass rate across:

- Phase 5 behavioral integrity tests
- Phase 6 context + intelligence route tests

Meaning:

Phase 7 did not break prior phases.

B. Phase 7 Engine Unit Tests

Dedicated test suites confirm:

- correct data access usage,
- correct context builder invocation,
- correct orchestrator call count,

- correct passthrough of SessionResponse.

C. Phase 7 Workflow Tests

Cross-engine integration tests confirm:

- Engines can run sequentially,
- Shared infrastructure is intact,
- State boundaries not violated.

D. End-to-End Smoke Validation

Phase 7 smoke script in /scripts correctly hits all four new endpoints successfully.

4. Code Health + Integrity Requirements

Structural

- Single intelligence tree
- No nested duplicate directory
- Single correct long_horizon.py
- No drift or fragmentation

Dependency Safety

- No engine violates orchestrator boundary
- No builder requires external runtime mutation
- All new tests aligned against real boundaries

Execution Integrity

- All engines execute without raising `TypeError`s
 - All context builders produce JSON-serializable outputs
-

5. Documentation

Two formal internal docs now exist:

- **Phase 7 Completion**
- **Phase 7 → Phase 8 Handoff**

Both accurately reflect the *actual* implementation and state of the codebase, not theoretical intent.

6. Final Status

Phase 7 Completion Verdict:

COMPLETE & VALIDATED

All architectural, functional, testing, and backward compatibility conditions have been met.

The codebase is stable and ready for Phase 8 development.
