

Phase 6 → Phase 7 Handoff Document

This document prepares the Phase-7 development environment and personnel by describing what Phase 6 provides and how Phase 7 should build on it. Phase 6 Delivered: - Fully modular intelligence engines (review, consistency, Q&A, commentary). - An orchestrator wrapper that unifies all intelligence calls. - Deterministic context layers for thesis, multi-thesis, RAG-based intuition, and commentary. - Fully tested API layer for intelligence. Phase 6 Exposed Contracts for Phase 7: 1. OrchestratorClient - Phase 7 should treat this wrapper as the sole gateway to the session pipeline. - Stable interface: `run_standard`, `run_analyst`, `run_concise`. 2. ContextBuilder Contracts - Guaranteed stable output structures for: `* thesis_context` `* multi_thesis_context` `* intuition_context` `* daily/weekly_context` 3. Intelligence Engine Contracts - Each engine returns `SessionResponse`. - Inputs are simple Python types. - No DB/I/O assumptions beyond `DataAccess` abstractions. 4. Intelligence API Surface - Five Phase-6 endpoints form the public interface for higher-level systems. - Phase 7 must preserve these routes and their response shapes. Phase 7 Objectives: - Introduce long-horizon reasoning features. - Expand intelligence engines to incorporate strategy recommendations. - Add external data sources for macro context. - Add portfolio-diagnostics engine. - Integrate more complex narrative-coherence systems. - Extend testing to include multi-engine workflows and cross-layer invariants. Guiding Principles for Phase 7: - Do not break `OrchestratorClient` boundaries. - Maintain structured context formation. - Avoid embedding logic inside routes. - Expand the intelligence layer without modifying Phase-5 internals. - Ensure all new engines follow the same `SessionResponse` pattern. This document provides the forward-looking blueprint to ensure Phase 7 builds cleanly atop the stabilized Phase-6 intelligence layer.